

Elements of Design for Containers and Solutions in the LinBox library

Brice Boyer¹, Jean-Guillaume Dumas², Pascal Giorgi³, Clément Pernet⁴, and
B. David Saunders⁵

¹ Department of Mathematics, North Carolina State University, USA[†]
`bbboyer@ncsu.edu`.

² Laboratoire J. Kuntzmann, Université de Grenoble. France[‡]
`Jean-Guillaume.Dumas@imag.fr`.

³ LIRMM, CNRS, Université Montpellier 2, France[‡]
`pascal.giorgi@lirmm.fr`.

⁴ Laboratoire LIG, Université de Grenoble et INRIA, France[‡]
`clement.pernet@imag.fr`.

⁵ University of Delaware, Computer and Information Science Department, USA
`saunders@udel.edu`.

Abstract. We develop in this paper design techniques used in the C++ exact linear algebra library LinBox. They are intended to make the library safer and easier to use, while keeping it generic and efficient.

First, we review the new simplified structure of the containers, based on our *founding scope allocation* model. Namely, vectors and matrix containers are all templated by a field and a storage type. Matrix interfaces all agree with the same minimal blackbox interface. This allows e.g. for a unification of our dense and sparse matrices, as well as a clearer model for matrices and submatrices. We explain the design choices and their impact on coding. We will describe several of the new containers, especially our sparse and dense matrices storages as well as their **apply** (*blackbox*) method and compare to previous implementations.

Then we present a variation of the *strategy* design pattern that is comprised of a controller–plugin system: the controller (solution) chooses among plug-ins (algorithms) and the plug-ins always call back the solution so a new choice can be made by the controller. We give examples using the solution `mul`, and generalise this design pattern to the library. We also show performance comparisons with former LinBox versions.

Finally we present a benchmark architecture that serves two purposes. The first one consists in providing the user with an easy way to produce graphs using C++. The second goal is to create a framework for automatically tuning the library (determine thresholds, choose algorithms) and provide a regression testing scheme.

Keywords: LinBox, design pattern, solutions and containers, benchmarking

[†] This material is based on work supported in part by the National Science Foundation under Grant CCF-1115772 (Kaltofen)

[‡] This material is based on work supported in part by the Agence Nationale pour la Recherche under Grant ANR-11-BS02-013 HPAC (Dumas, Giorgi, Pernet).

1 Introduction

This article follows several papers and memoirs on the LinBox⁶ (*cf.* [10,14,2,6,7]) and builds upon them.

LinBox is a C++ template library for fast and exact linear algebra. It is designed with genericity and efficiency in mind. The LinBox library is under constant evolution, driven by new problems and algorithms, by new computing paradigms, new compilers and architectures. This poses many new challenges. To address this changes, we are incrementally updating the *design* of the library towards a 2.0 release.

Let’s start from a basic consideration: we show in the Table 1 the increase in the size⁷ of LinBox and its dependancies in terms of “lines of code”. This

LinBox	1.0.0 ^{†‡}	1.1.0 ^{†‡}	1.1.6 [‡]	1.1.7 [‡]	1.2.0	1.2.2	1.3.0	1.4.0
loc (×1000)	77.3	85.8	93.5	103	108	109	112	135
FFLAS–FFPACK	n/a	n/a	n/a	1.3.3	1.4.0	1.4.3	1.5.0	1.8.0
loc	—	—	—	11.6	23.9	25.2	25.5	32.1
Givaro	n/a	n/a	3.2.16	3.3.3	3.4.3	3.5.0	3.6.0	3.8.0
loc	—	—	30.8	33.6	39.4	41.1	41.4	42.8
total	77.3	85.8	124	137	171	175	179	210

Table 1: Evolution of the number of lines of code (loc, in thousands) in LinBox, FFLAS–FFPACK and Givaro ([†]contains Givaro, [‡]contains FFLAS–FFPACK).

increase affects the library in several ways. First, it demands a stricter developpement model, and we are going to list some techniques we used. For instance, we have transformed FFLAS–FFPACK⁸ (*cf.* [8]) into a new stand-alone header library, resulting in more visibility for the FFLAS–FFPACK project (Singular ?) but also in a better structuration an maintainability of the library, focusing the developpement areas more precisely. Also, a larger template library is harder to manage, there is more difficulty to trace, debug and write new code: techniques employed for easier developpement include reducing compile times, enforcing stricter warnings and checks, supporting for more compilers and more architectures, simplifying and automatising version number changes, automatising memory leak checks, setting up buildbots to check the code frequently,...

But this increase also forces the library to be more user friendly. For instance, we have: Developped an `auto-install.sh` script that installs automatically the latest stable or developpement versions of the trio; Facilitated the discovery of the BLAS/LAPACK libraries; Simplified and sped up the checking process while covering more of the library (\triangle dave ?); Added comprehensive benchmarking tools,...

⁶ See <http://www.linalg.org>.

⁷ Using `sloccount`, available at <http://sourceforge.net/projects/sloccount/>.

⁸ See <http://www.linalg.org/projects/fflas-ffpack/>.

Developping generic and high-performance libraries is difficult. We can find a large litterature on coding standards and software design references in (*cf.* [1,9,13,12,11]), and many internet sources and a lot of experience acquired by/from free software projects.

We are going to describe the advancement in the design of LinBox in the next three sections. We will first describe the new *container* framework in Section 2, then improve the *matrix multiplication* algorithms in Section 3 by contributing special purpose matrix multiplication plugings, and finally present the new *benchmark/optimisation* architecture (Section 4). \triangle develop this § more later

2 Containers architecture

LinBox is mainly conceived around the RAII concept with re-entrant function (Resource Acquisition Is Initialisation), introduced by [12]. We also follow the founding scope allocation model (or *mother model*) from [7] which ensures that the memory used by objects is allocated in the constructor and freed only at its destruction. The gestion of the memory allocated by an object is then exlusively reserved to it.

LinBox essentially uses matrix and vectors over fields as data objects (containers). The fragmentation of the containers into various matrices and blackboxes needed to be addressed and simplified. The many different matrix and vector types with different interfaces needed to be reduced into only two (possibly essentially one in the future) containers: **Matrix** and **Vector**.

2.1 General Vectors and Matrices

Firstly, in order to allow operations on its elements, a container is parametrized by a field, not the element type; this is also more general. The storage type is given by another template parameter that can default to *e.g.* dense BLAS type matrices, *cf.* the Listing 1.1.

```

5  template< class _Field, class _Storage = denseDefault >
    class Matrix ;

    template< class _Field, class _Storage = denseDefault >
    class Vector ;

```

Listing 1.1: Matrix and vector classes in LinBox.

Interface of the containers. In the mother model, we need types that own or and types that share some memory. The **SubMatrix** and **SubVector** types share the memory while **Matrix** and **Vector** own it. The common interface to all matrices is the **BlackBox** interface described in Listing 1.2.

```

template< ... >
class Matrix {
public:
    /* constructors */
    // no empty constructor
    // with a field, rows, columns, other matrix\dots

    /* y <- A.x or y <- A^T.x */
    template<class _in, class _out>
    _out& apply(_out &y, const _in &x) const;
    template<class _in, class _out>
    _out& applyTranspose(_out &y, const _in &x) const;

    /* operator rebind */
    template<typename _Tp1>
    struct rebind ;

    /* dimensions */
    size_t rowdim() const;
    size_t coldim() const;

    /* field */
    const Field& field() const;

    /* conversions */
    // resize
    // export/import to default types

    /* building the matrix */
    // read from MM/stdout/...
    // write to MM/stdout/...
    // setEntry/init()/finalise()/optimize()
protected:
    /* internals */
    ...
};

```

Listing 1.2: Interface of an BlackBox

Remark 1. Some functions are not part of the default interface :

- **refEntry** that retrieves a reference to an entry may be difficult to implement or inefficient (compressed fields, sparse matrices)
- **getEntry** may be specialized, in all cases, there is a solution for this operation (can always be implemented from **imply**, *cf.* later.
- **clearEntry** can be used to zero out an entry, especially for a sparse matrix, if this is allowed (possibly not for structured matrices).
- iterators may be difficult to implement (but a lot of code relies on them...). Do we want only **const** iterators ?

The function `setEntry(...)` can be used to populate/grow the matrix (from some `init()` until a `finish()` is emitted). The function `setEntry` can be (very) costly (for some sparse formats for instance) (Dave ?)

Input/Output. Our matrix all read and write from MatrixMarket format (ref, link). Adding extra comments ? (for instance the init function in GF(q) needs a polynomial...)

2.2 The apply method

We will discuss the `apply` function in `vec/mat`

The `apply` method (left or right) is arguably the most important feature in the matrix interface and the LinBox library. It performs what a linear application is defined for: apply to a vector (and by extension a block of vectors, *i.e.* a matrix).

Proposed interface :

```

// y = A.x
template< class _In, class _Out >
_Out& apply(_Out &y, const _In& x, enum Side) ;

5 // y = alpha y + beta A.x
template< class _In, class _Out >
_Out& applyAcc(_Out &y, const Element& alpha, const _In& x, const
      Element& beta, enum Side) ;

```

Listing 1.3: Proposed apply.

The `apply` method should be using a `mul` solution as in the Listing 1.4.

```

template<class _outVector, class _inVector>
OutVector &apply (_outVector &y, const _inVector &x) const
{
    /* return _MD.vectorMul (y, *this, x); */
5    return blas3::mul(y, *this, x); // selects the best mul
        algorithm
}

```

Listing 1.4: Implantation d'apply sur un vecteur par solution mul.

2.3 Sparse Matrix

Sparse matrices are usually problematic because the notion of *sparsity* is too general and the matrices we use are usually very specific: the algorithms have to adapt to the shape of the sparse matrices. Getting the best performance for an

sparse matrix as an **BlackBox** is not an easy task (*cf.* [3]). There is a huge litterature on sparse matrix formats, some of which are becoming standard. Numerical analysis brings various shapes for sparse matrices and numerous algorithms and routines. Just like the BLAS numerical routines, we would like to take advantage of existing high performance libraries. They are however not very widespread, for instance sparse blas in intel mkl (only).

△metis, sparse suite, zero-one matrices

Matrix market matrices must be supported (we have a convertor in ffsp-mvcpu, I should commit it)

SparseMatrix are Matrix with specified storage : COO, CSR, CSC, ELL, ELL_R,..., and more to come. (including 0-1 based)

```

5      namespace matrixStorage {
          struct sparse {};
          struct COO : public sparse {};
          struct CSR : public sparse {};
          struct ELL : public sparse {};
          struct ELLR : public sparse {};
          struct HYB : public sparse {};
          ...
10    };

    template<class _Field >
    class Matrix<_Field, matrixStorage::COO> {
        //specialisation for COO
    };

```

Listing 1.5: Standard sparse matrix formats.

HYB format is a specialisation for ELL(_R)+COO/CSR. no easy setEntry (in the COO/CSR) part ? a function `optimize()`. Could use Metis too. litterature for other HYB.

We use MM reader.

Added to the interface is a convert method to/from CSR (default).

We can adapt the header to suit our needs. In particular write matrix in CSR fashion (saving roughly 1/3 space over COO)

XXX timing new matrices, example rank.

Other Matrices. All other matrices (including the special case of Permutations) : the same Structured matrices, add, sub, stacked,...

3 Improving LinBox matrix multiplication

Efficient matrix multiplication is key to LinBox library.

3.1 Plugin

We propose the following design pattern (the closest design pattern to our knowledge is the *strategy* one, see also [5, Fig 2.]). The main advantage of this design pattern is that the modules always call back the controller so that the best choice is always chosen. Besides modules can be easily added as *plug-ins*. An analogy can be drawn with dynamic systems—once the controller sends a correction to the system, it receives back a new measure that allows for a new correction.




Fig. 1: Controller–Modules design pattern

For instance, we can write the standard cascade algorithms in that model:

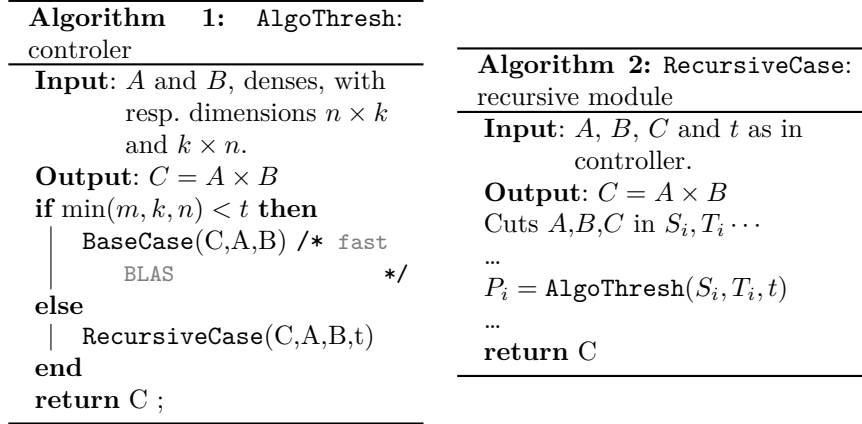


Fig. 2: Conception of a recursive controlled algorithm

This method allows for the reuse of modules and ensures efficiency. It is then possible to adapt to the architecture, the available modules, the resources. The only limitation is that the choice of the module should be done fast. △timing old fgemm/plugin fgemm with no noticeable change ?

3.2 New algorithms/infrastructure

We introduce now several new algorithms that improve on matrix multiplication in various ways: reducing memory consumption, introducing new efficient algorithms, using graphics capabilities, generalizing the BLAS to integer routines.

New algorithms: low memory `ffgem` in FFLAS uses the classic schedules for the multiplication and the product with accumulation (*cf.* [4]), but we also implement the lower memory routines therein.

The difficulty consists in using the part of the memory contained in a sub-matrix of the original matrix. It is two-fold. – First we use some part of a memory that has already been allocated to the input matrices, therefore we cannot free and reallocate part of it. – Second, several of these algorithms are meant for square matrices and rectangular sub-matrices will just not be enough. For instance,

△table comparing speeds

New algorithms: Bini

integer blas △pascal

OpenCL △dave

Using conversions - using flint for integer matmul is faster, even with conversion. Need better CRA implementation.

- implementation of Toom-Cook for GF(q)

- when does spmv choose to optimise ?

4 Benchmarking

Benchmarking was introduced in LinBox for several reasons. It would give the user a user-friendly way for producing nice graph with no necessary knowledge of the graphing library `gnuplot`⁹ or provide LinBox website with automatic new tables/graphs. It would be used for regression testing. It can be used for selecting default method, threshold. (cite atlas and?)

XXX needs introductory bench references

XXX BTL¹⁰/eigen △dave benchmark formats

4.1 Graph/Table creation

XXX data/plot structure of abstraction (update code 6.1 with new plot structure)

⁹ <http://www.gnuplot.info/>

¹⁰ <http://projects.opencascade.org/btl/>

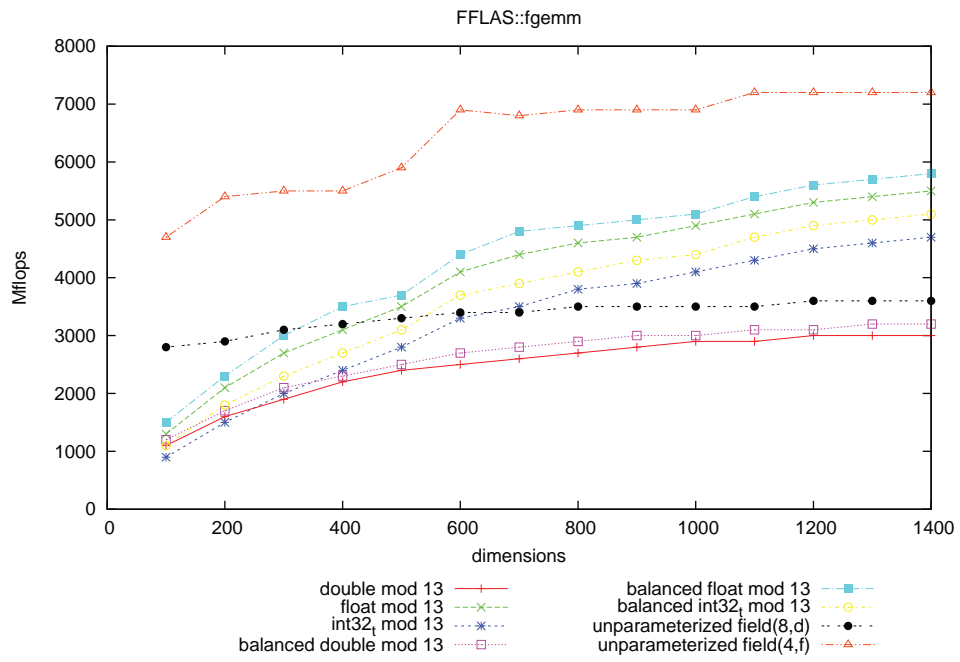


Fig. 3: Example of benchmark: fgemm.

XXX Time spent on each data is limited (will not start execution if fit (linear least squares) forecasts 'too long').
 XXX Adapts to the environment.

4.2 Regression Testing

XXX Tests are supposed to be comparable: regression testing.
 XXX howto.

4.3 Method Selecting

XXX Default are provided, method can be selected via a benchmark (cf `wino_threshold`)
 XXX howto

References

1. A. Alexandrescu. *Modern C++ design: generic programming and design patterns applied*. C++ in-depth series. Addison-Wesley, 2001.
2. B. Boyer. *Multiplication matricielle efficace et conception logicielle pour la bibliothèque de calcul exact LINBOX*. PhD thesis, Université de Grenoble, June 2012.
3. B. Boyer, J.-G. Dumas, and P. Giorgi. Exact sparse matrix-vector multiplication on GPU's and multicore architectures. In *Proceedings of the 4th International Workshop on Parallel and Symbolic Computation*, PASCO '10, pages 80–88, New York, NY, USA, 2010. ACM.
4. B. Boyer, J.-G. Dumas, C. Pernet, and W. Zhou. Memory efficient scheduling of Strassen-Winograd's matrix multiplication algorithm. In *Proceedings of the 2009 international symposium on Symbolic and algebraic computation*, ISSAC '09, pages 55–62, New York, NY, USA, 2009. ACM.
5. V.-D. Cung, V. Danjean, J.-G. Dumas, T. Gautier, G. Huard, B. Raffin, C. Rapine, J.-L. Roch, and D. Trystram. Adaptive and hybrid algorithms: classification and illustration on triangular system solving. In J.-G. Dumas, editor, *Proceedings of Transgressive Computing 2006, Granada, España*, Apr. 2006.
6. J.-G. Dumas, T. Gautier, M. Giesbrecht, P. Giorgi, B. Hovinen, E. Kaltoven, B. D. Saunders, W. J. Turner, and G. Villard. LINBOX: A generic library for exact linear algebra. In *Proceedings of the 2002 International Congress of Mathematical Software, Beijing, China*. World Scientific Pub, Aug. 2002.
7. J.-G. Dumas, T. Gautier, C. Pernet, and B. D. Saunders. LinBox founding scope allocation, parallel building blocks, and separate compilation. In K. Fukuda, J. Van Der Hoeven, and M. Joswig, editors, *Proceedings of the Third International Congress Conference on Mathematical Software*, volume 6327 of *ICMS'10*, pages 77–83, Berlin, Heidelberg, Sept. 2010. Springer-Verlag.
8. J.-G. Dumas, P. Giorgi, and C. Pernet. Dense linear algebra over word-size prime fields: the FFLAS and FFPACK packages. *ACM Trans. Math. Softw.*, 35(3):1–42, 2008.
9. E. Gamma. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley Professional Computing Series. Addison-Wesley, 1995.

10. P. Giorgi. *Arithmétique et algorithmique en algèbre linéaire exacte pour la bibliothèque LINBOX*. PhD thesis, École normale supérieure de Lyon, Dec. 2004.
11. D. Gregor, J. Järvi, M. Kulkarni, A. Lumsdaine, D. Musser, and S. Schupp. Generic programming and high-performance libraries. *International Journal of Parallel Programming*, 33:145–164, 2005. 10.1007/s10766-005-3580-8.
12. B. Stroustrup. *The design and evolution of C++*. Programming languages/C++. Addison-Wesley, 1994.
13. H. Sutter and A. Alexandrescu. *C++ Coding Standards: 101 Rules, Guidelines, And Best Practices*. The C++ In-Depth Series. Addison-Wesley, 2005.
14. W. J. Turner. *Blackbox linear algebra with the LINBOX library*. PhD thesis, North Carolina State University, May 2002.